

Journalists' Challenges with Digital Content Provenance and Trust:

How Semantic Knowledge Graphs Enable Traceability and Context

by Dinis Cruz and ChatGPT Deep Research, 24 March 2025

Digital content is being produced at an unprecedented scale, largely driven by AI technologies. Generative AI can churn out articles, images, and videos in seconds, contributing to a veritable flood of content online. *“For good and bad, AI-written content will flood the internet, an amalgamation of the work of millions of human writers and journalists who came before it,”* as one media expert put it ([The AI content flood » Nieman Journalism Lab](#)). In this environment of **digital abundance**, the competition for audience attention is fierce, and the **integrity** of information is under strain. Audiences are increasingly asking: *“Can I trust what I’m reading?”*

For journalists, whose work relies on credibility, this question is critical. Trust has in fact become *“the internet’s most valuable asset”* in an age where content is abundant but attention and confidence are scarce ([After Attention: Trust in the Age of Digital Abundance](#)). Establishing **provenance** – a clear record of where information comes from – is emerging as a key strategy for rebuilding trust. This article delves into why provenance and trust matter more than ever, the challenges posed by AI-driven content commoditisation, and how **Semantic Knowledge Graphs** offer a solution by infusing traceability, context, and credibility into news production.

The Commoditisation of Content in the Age of AI

The rise of generative AI has effectively **commoditised digital content**. What does this mean? Essentially, content creation has become cheap, automated, and mass-produced. Today we stand at a juncture marked by *“the commoditisation of digital content driven by Generative AI (GenAI)”* ([Generative AI Media Production | Deloitte US](#)). News articles, blog posts, and even deepfake videos can be generated with minimal human effort. The result is **mass content generation**: an endless stream of articles and posts, often of varying quality. When content becomes so plentiful, it starts to lose

its distinctive value – a well-researched investigative piece can appear, at a glance, in the same feed as a hastily generated AI summary. This oversupply dilutes the perceived value of each piece of content.

A direct consequence of this flood is **decreasing public trust**. Audiences know that not everything online is written by a responsible human journalist; it could be machine-generated, unverified text. People worry about misinformation and “fake news” slipping in. If any website can publish dozens of AI-written stories per day, how can one tell which stories went through careful journalistic vetting? This uncertainty chips away at trust in digital media broadly. In short, the **integrity signal** that traditionally came with professional journalism is at risk of being drowned out by the noise of commoditised content.

Dilution of value is another challenge. Quality journalism – with on-the-ground reporting, fact-checking, and expert analysis – is costly to produce. But when readers are flooded with free auto-generated content, they may become less willing to pay for or even seek out quality news. The temptation to treat news as an interchangeable commodity grows. For journalists and news organisations, this is a call to action: they must find new ways to **demonstrate value and trustworthiness** to set their work apart from the glut of low-effort content.

Why Provenance and Trust Matter for Journalists

Trust is the bedrock of journalism ([Trust is the bedrock of journalism. How can news and media companies... | Content Authenticity Initiative](#)). Without trust, even the most important story will fall flat with the public. That’s why journalists have long upheld standards like verifying sources, attributing information, and issuing corrections – all efforts to maintain credibility. In the digital era, **provenance** has become a vital concept in preserving that trust. Provenance refers to the history or origin of a piece of content – essentially an audit trail that can show where information came from and how it has been handled or changed over time ([Watermarks are Just One of Many Tools Needed for Effective Use of AI in News - Center for News, Technology & Innovation](#)). In the art world, provenance might tell you the chain of owners of a painting. In journalism, provenance means having documentation of sources, evidence, and changes made during editing.

Importantly, provenance is even more crucial for journalists and content creators than for the general audience. It is part of the **internal fabric of accountability** in news production. As one technology leader noted, *“the value of provenance is probably more for the world of journalism than it is for the world of consumers.”* Journalists need to know the source of information they rely on, and editors need to verify that reporting is solid. In an age of potential AI manipulation, having a **manifest trail of information** (who said what, when, and where) is indispensable.

When provenance is lacking or opaque, misinformation can spread more easily. A tidbit from an unverified blog might get repeated by larger outlets without clear attribution, eventually snowballing into a widely believed falsehood because nobody checked the original source. We have seen how quickly **rumors and deepfakes** can ricochet through social media, often because the origin is obscured or forgotten. For journalists, establishing provenance isn't just an academic exercise – it's a defense against being duped by bad information and a way to **uphold the truth** in their storytelling.

In an **increasingly confusing media landscape**, bringing clarity through transparency is key. Many news organisations and tech companies are now exploring tools (like digital watermarks or cryptographic content signatures) to authenticate content. But a complementary approach lies in structuring the content's information itself in a way that reinforces trust. This is where **Semantic Knowledge Graphs** come into play as a promising solution.

Challenges Posed by AI-Driven Content Proliferation

Before exploring the solution, let's summarise the challenges that modern newsrooms face due to AI-driven content proliferation:

- **Mass AI Generation & Misinformation:** Generative AI can produce a deluge of text that looks credible but may contain errors or falsehoods. Mis/disinformation actors can also use AI to generate fake news at scale. The volume makes it hard for readers (and even journalists) to separate fact from fiction when every piece looks similar.
- **Erosion of Audience Trust:** People have grown wary as they realise that not all content is human-authored or fact-checked. Credible outlets risk losing trust by association if audiences feel "anything online could be fake." Establishing trust now requires extra effort in showing authenticity and accuracy.
- **Dilution of Journalistic Value:** With AI automating content, there's pressure on news organisations to produce more with less, potentially sacrificing quality. It also becomes harder for quality journalism to **stand out** when low-quality content is plentiful and often free. The economic model of journalism is challenged when content is treated as a cheap commodity.
- **Verification Bottleneck:** The duty to fact-check and verify becomes both more important and more difficult with so much content. Newsroom resources are strained trying to keep up with debunking false claims that proliferate online.

These challenges underscore why simply *producing* truthful journalism isn't enough; journalists must also actively **signal and prove** their integrity and the provenance of their information. This sets the stage for adopting new tools that can bolster trust.

Semantic Knowledge Graphs: A Solution for Traceability and Context

One of the most promising tools to address these issues is the **Semantic Knowledge Graph**. In essence, a semantic knowledge graph is a structured representation of information that captures entities (people, places, things) and the relationships between them. Instead of treating a news article as just text, a knowledge graph turns it into **data**: who is involved, what happened, where and when it happened, and how those elements connect. In practice, this often means using AI to analyse an article and output a set of triples or nodes: for example, (Person A) -- [participated in] --> (Event X) .

To see this in action, consider the approach of the MyFeeds.ai MVP, which uses semantic knowledge graphs under the hood. Each incoming article (in their case, cybersecurity news) is analysed and **converted into a knowledge graph of the article's key entities and their relationships** ([Establishing Provenance and Deterministic Behaviour in an LLM-Powered News Feed \(first MyFeeds.ai MVP\)](#)). Essentially, the system is extracting the *who*, *what*, and *how* of the story and encoding those facts into a graph data structure. Each article yields its own graph (represented in JSON) capturing the essential concepts in machine-readable form.

Even with just an article's title and summary, a semantic graph can reveal a **dense web of interconnected entities**, as shown in the figure below. This illustrates how much structured information can be derived from raw text, far beyond a simple keyword list. The graph's nodes might include people, organisations, locations, and topics mentioned, and edges describe the relations (e.g., "works for", "located in", "spoke about", etc.). Such a graph effectively provides **context at a glance**: it shows the broader network of relationships that the article touches on, placing a news story within a larger framework of knowledge.

By structuring news content this way, we gain several advantages: **traceability, context, and credibility**. In the following sections, we break down how semantic knowledge graphs specifically address the trust and provenance challenges in digital journalism.

Traceability: Building a Provenance Trail for News

A core benefit of using knowledge graphs in news is the creation of a **traceable provenance trail**. Because a knowledge graph is essentially a database of facts extracted from the article, it can also include metadata about *sources* and *origins* of those facts. Every node or edge in the graph can carry an

attribute pointing to *where that piece of information came from*. For example, a node representing a quoted statement could link to the source (say, a press release or an interview recording), and an edge might indicate it was *quoted by* the article.

In the MyFeeds example, the system not only generates graphs but also uses them to explain its output. Each stage of content processing produces structured data, which means the system can answer **“Why am I seeing this piece of news?”** with concrete evidence. The chain of outputs forms a provenance trail that shows why a particular article was recommended_ to a reader. For instance, if a personalised news feed surfaces an article about “GraphQL” for a CTO persona, the system can point to the fact that the article mentions GraphQL, and the user’s persona profile lists GraphQL as an interest. This step-by-step record makes the process transparent and auditable. Each connection is explicit and backed by data, turning what could be a black-box recommendation into an **open, explainable decision**.

The same principle applies to editorial decisions in a newsroom. Imagine an investigative article that draws on five sources: a knowledge graph of that article could explicitly link every major claim to one of those source documents. The journalist (or an automated system) can then trace each fact: *“This statistic came from Source A (with a hyperlink or reference), this quote came from Source B,”* and so on. If any detail is challenged, the provenance trail is readily available to backtrack and verify. This is essentially automating a part of the journalist’s notebook – capturing the **attribution for every piece of information** in a structured way.

Another area where traceability is vital is in **fact-checking and verification workflows**. Journalistic Knowledge Graph Platforms (JKPs) have begun to integrate verification tasks into their functionality. Such platforms *“support news professionals with verification tasks like fact-checking and provenance tracking”* ([Supporting Newsrooms with Journalistic Knowledge Graph Platforms: Current State and Future Directions](#)). By automating and logging these steps, knowledge graphs ensure that verification isn’t an afterthought but a built-in feature of content creation. They can flag, for example, if a fact in the story doesn’t align with known data in a reference knowledge base, prompting a journalist to double-check. All of this contributes to **content credibility** – when challenged, a news outlet can produce a clear trail of how information was gathered and vetted.

Traceability via knowledge graphs thus acts as a **“nutritional label”** for news content, to borrow an analogy used in industry discussions. Just like a food label tells you the origin and ingredients, a provenance-rich knowledge graph tells you the sources and components of a news story. Over time, if such labels become common, audiences may come to expect and trust content that carries transparent provenance, and conversely be skeptical of content that does not. It creates a feedback loop that rewards trustworthy journalism.

Providing Context and Capturing Editorial Intent

Beyond provenance, knowledge graphs excel at preserving **context** – the subtle connections and background that give meaning to facts. For journalists, context is king: a raw fact is less useful without understanding its significance or relation to other facts. Semantic graphs inherently capture relationships that provide this context. For example, a knowledge graph for a political news story wouldn't just list the politicians involved; it might show that *Politician X is part of Organisation Y*, or *Policy Z was previously proposed in Year T*. These connections help both journalists and readers see the *bigger picture* around a news event.

Consider how a newsroom might use such a graph. A reporter writing about a new climate agreement could quickly query the knowledge graph for related treaties, key players, and past events on climate policy. The graph, populated from archives and databases, effectively surfaces *the context that might not be in the immediate story but is crucial for understanding it*. This addresses a common issue in digital media: articles often appear in isolation, and readers may lack background. A knowledge graph can power sidebars or links like “*Related: Previous Agreement*” or interactive visuals showing how this event ties into a timeline or network of events. In other words, it helps maintain the **continuity of news** – connecting today's story with yesterday's developments and tomorrow's ramifications.

Editorial intent can also be captured or at least reflected in knowledge graphs. Editorial intent refers to the purpose or angle behind a story – what the journalist or news outlet aims to convey. While intent itself might be abstract, a knowledge graph can demonstrate *which aspects of a story were highlighted*. For instance, if an editor decides that an article should focus on the cybersecurity angle of a new technology (because their audience cares about that), the final article's knowledge graph might have a rich set of nodes and relationships in the cybersecurity domain (threats, breaches, tools mentioned) relative to other possible angles. This is a form of semantic fingerprint of the editorial choices made.

The MyFeeds MVP provides a concrete example of encoding something akin to editorial intent through **personalisation graphs**. It creates a “persona” knowledge graph representing a reader's interests (say, a CTO cares about cloud infrastructure and data breaches). Then it compares an article's graph to the persona graph to find relevant overlaps. The overlap determines what to emphasise in a summary, effectively tailoring the editorial output to that persona. In a broader sense, this shows how graphs can align content with the intended audience or purpose. If we extend this idea, a newsroom could have multiple schema or subgraphs to represent different editorial frames (e.g. economic impact, human interest, political angle) and ensure that the content includes nodes from those frames according to the editorial plan.

Moreover, knowledge graphs support **narrative tracking**. Journalists often follow story threads over weeks or months. A knowledge graph can link stories by events, people, and themes, so one can trace how a narrative evolves. This is very useful for editors planning follow-ups or investigative reporters looking for connections. It also helps prevent context loss: when a new reporter takes over a beat, the established knowledge graph is a treasure trove of what's been covered and how things relate.

Ultimately, by making the context explicit, semantic graphs ensure that journalism doesn't happen in a vacuum. They help maintain the richness of a story's background and the clarity of its focus, both of which contribute to **audience trust**. Readers are more likely to trust reporting that seems well-contextualised and comprehensive, and journalists can do their jobs better when they have a map of the knowledge terrain around their story.

Source Validation and Fact-Checking with Knowledge Graphs

Every journalist knows the mantra: “**Check your sources**.” In the AI era, source validation is both crucial and challenging. Knowledge graphs can become powerful allies in this respect by incorporating source metadata and even fact-checking logic directly into the news data model.

A semantic knowledge graph for news can include not just the *content* of sources, but also information about the **sources themselves**. For example, each source (be it a document, website, or person quoted) can be a node in the graph with attributes like author, publisher, date, and even a credibility rating. Research into news credibility is already producing “source credibility graphs” that rate outlets or authors on trustworthiness. A newsroom could integrate such data, so that if an article pulls information from a source, the graph carries along the credibility meta-information. This allows both journalists and algorithms to weigh information accordingly – e.g., a claim from a highly credible source might be treated with more default trust than one from an unknown blog.

Moreover, by having facts and claims as first-class elements in the knowledge graph, it becomes easier to **fact-check** them. Automated systems or editors can cross-reference claims against established knowledge bases. In some experimental systems, knowledge graphs are used in *fake news detection* by comparing claims to known factual triples. If a new statement contradicts what's in a trusted knowledge graph (like Wikidata or a curated database), it's a red flag that warrants further scrutiny. Conversely, if the knowledge graph confirms a claim (or the source provides evidence in the graph), that fact is on firmer ground.

Journalistic Knowledge Graph Platforms often incorporate these verification features. They “*facilitate access to previous work*” and knowledge bases, and “*support fact-checking (and) provenance*” steps as part of content creation. This means that when a journalist is writing a story and adds a

statement, the platform could automatically fetch related facts or past fact-checks, showing, for example, that *“Claim X was rated false in a fact-check last year”* or *“Statistic Y was reported in these sources with a different figure.”* Having this real-time or built-in fact-checking drastically improves newsroom efficiency and reliability. It’s like having a smart assistant that constantly checks the draft against a repository of truth.

When it comes to **differentiating trustworthy information**, validated sources in a knowledge graph shine. Each piece of content can come with a kind of pedigree: not just *what* the information is, but *who vouched for it*. For instance, a graph might show that a medical claim in an article is linked to a World Health Organisation report (a strong source), or that a political quote was verified by a known fact-checking organisation. This builds a web of trust around the content.

It’s also worth noting that the graph approach can capture the **structure of fact-checking** itself. Fact-check articles could be represented in the graph with links to the original claims they checked and the evidence used. Over time, this creates a lattice of claims and verifications, which both human journalists and AI tools can use to quickly assess new statements. If a new article repeats a debunked claim, the system can flag it by tracing a path in the graph to a fact-check node that marked it false. Such proactive alerting is invaluable when misinformation can spread so quickly.

By validating sources and facilitating fact-checking, semantic knowledge graphs directly tackle the credibility crisis. They help ensure that what gets published has been vetted against reliable information stores. And equally important, they **show the work** – making the validation process visible and traceable. For news audiences, knowing that a story’s facts are backed by an accessible web of sources and verifications can significantly bolster confidence. It’s the antidote to the opaque “trust me” approach; instead, the attitude becomes *“see for yourself – here’s how you know this is true.”*

Human vs. AI Content: A New Differentiator

In a world awash with auto-generated text, having strong provenance and semantic richness can be the key to **differentiating trustworthy, human-crafted journalism from AI-generated or low-trust content**. As standards and tools for provenance gain adoption, we may reach a pivotal point: *“fact-based content gets promoted while ... inauthentic material or material of unknown provenance is lessened”*, as one industry report suggested. In other words, content that carries transparency about its creation and sources will be elevated, and content that hides its origins will be deprioritised. The goal is that *“only bad actors are the ones not willing to use a system like this”*. If that threshold is reached, it becomes much easier for both platforms and readers to spot and sideline dubious content. Lack of provenance will itself be a warning sign.

Semantic knowledge graphs play a big role in achieving this vision. Because they encapsulate so much contextual and source information, they can serve as a **credibility signature** for content. A well-developed news knowledge graph behind an article indicates that the article is the product of an information-rich process: sources were tracked, context was considered, facts were connected to evidence. It is unlikely that a purely AI-generated “slop” piece (to use a recent term for low-quality AI content) would come packaged with a detailed knowledge graph complete with source attributions and editorial context. Creating such a graph requires either diligent human oversight or advanced AI pipelines explicitly designed for transparency – either way, a level of effort that low-trust content farms probably won’t invest in.

Going forward, we might see news products where readers can **toggle a transparency view** – essentially peeking at the knowledge graph of an article to see its DNA. Journalistic content that has this interactive map will immediately signal its legitimacy. By contrast, if an article can’t show how it knows what it claims, readers may rightfully question, “Was this just made up by a machine?” Journalists can take pride in the fact that their work, when supported by such graphs, contains *depth* that algorithmic pastiche cannot match. It’s akin to the difference between a scholarly article with citations and a random blog post with no references – one clearly shows rigorous work, the other asks to be taken on faith.

This differentiation also has an algorithmic side: search engines and social media algorithms could leverage the presence of provenance data and knowledge graph completeness as a ranking signal. Just as sites with better security (HTTPS) or better mobile layouts got ranking boosts in the past, perhaps in the near future content with verifiable provenance and rich semantic markup will be favored. This creates a strong incentive for newsrooms to adopt knowledge graph approaches – not only for ethical reasons but because it will directly impact reach and discoverability.

In summary, semantic knowledge graphs help enforce a new **integrity standard** in digital media. They equip journalism with a sort of *digital watermark of trust*, one that is far more informative than any visible badge or logo. It’s a machine-readable and human-auditable imprint that says: “*This piece of content has been through a process of truthful assembly, here is the proof.*” By embracing this, journalists and reputable media can clearly distinguish themselves from the sea of AI-generated content that comes with no such assurances.

Key Takeaways

- **The AI Content Flood & Trust Crisis:** Generative AI has flooded the internet with content, effectively commoditising it. In this abundance, trust has become the scarcest resource, making provenance (traceable origin of information) critical for journalists and audiences alike. Quality journalism now competes with a vast amount of low-quality or automated material, underscoring the need to prove its credibility at every turn.

- **Provenance as a Pillar of Integrity:** Provenance provides an **audit trail** for content, showing who created it and where the information came from. This is especially important in journalism – reporters need robust source trails to maintain accuracy. Without provenance, misinformation can propagate unchecked. With it, journalists can uphold transparency and accountability, reinforcing trust in their reporting.
- **Semantic Knowledge Graphs – Traceability & Context:** Semantic knowledge graphs offer a powerful way to encode provenance and context. They transform articles into networks of facts, entities, and sources, creating a **step-by-step provenance trail** for each piece of information ([Establishing Provenance and Deterministic Behaviour in an LLM-Powered News Feed \(first MyFeeds.ai MVP\)](#)). This enables traceability (any claim can be traced back to a source) and rich context (readers and editors can see how each story fits into a larger web of events and facts). In practical terms, these graphs turn opaque news into **explainable, data-backed stories**.
- **Credibility through Verification and Intent Mapping:** Knowledge graphs can incorporate source credibility and fact-checking, automatically flagging dubious claims and highlighting verified information. They also capture editorial intent and emphasis by showing which aspects of a story are linked to key themes or reader interests. The result is content that is not only factual but also transparent about its angle and evidence. This helps audiences discern genuine journalism from AI-generated text with no accountability.
- **Differentiating Trustworthy Content in the AI Era:** Embracing provenance and semantic graphs will increasingly separate reputable journalism from low-trust content. As provenance standards catch on, content that lacks a transparent knowledge graph or source trail will stand out—for the wrong reasons. Meanwhile, journalism that provides this data will gain **competitive advantage in visibility and trust**. The aim is to reach a point where not using such transparency tools marks a publisher as a likely bad actor. In short, **semantic knowledge graphs could become the integrity infrastructure that allows truth to thrive even as AI continues to proliferate content**.

By focusing on provenance and leveraging semantic knowledge graphs, journalists and news organisations can turn the very forces that threaten to dilute media trust into opportunities. They can harness AI to add structure and transparency to their work, ensuring that in the deluge of digital content, the stories that matter – the truthful, context-rich, human-authored stories – are the ones that shine through most clearly.